

Lecture 20 - Low-Level Programming

Meng-Hsun Tsai
CSIE, NCKU


| | | | | | |
|-----|----|-----|------------|-------------|-----|
| & | | ~ | << | >> | ^ |
| ↙ | ↙ | ↙ | ↙ | ↘ | ↘ |
| AND | OR | NOT | Left shift | Right shift | XOR |

Introduction

- Some kinds of programs need to perform operations at the bit level:
 - **Systems programs** (including **compilers** and **operating systems**)
 - **Encryption** programs
 - **Graphics** programs
 - Programs for which **fast execution** and/or **efficient use of space is critical**

Bitwise Operators

- C provides **six *bitwise operators***, which operate on **integer data** at the **bit level**.
- **Two** of these operators perform **shift** operations.
- The other **four** perform **bitwise complement, bitwise *and*, bitwise exclusive *or*, and bitwise inclusive *or*** operations.



Bitwise Shift Operators

- The bitwise shift operators shift the bits in an integer to the left or right:
 - << left shift
 - >> right shift
- The **operands** for << and >> may be of any integer type (including `char`).
- The **integer promotions** are performed on both operands; the **result** has the type of the left operand after promotion.

Bitwise Shift Operators (cont.)

- The value of $i \ll j$ is the result when the bits in i are shifted left by j places.
 - For each bit that is “shifted off” the left end of i , a zero bit enters at the right.
- The value of $i \gg j$ is the result when i is shifted right by j places.
 - If i is of an unsigned type or if the value of i is nonnegative, zeros are added at the left as needed.
 - If i is negative, the result is implementation-defined.

Bitwise Shift Operators (cont.)

- Examples illustrating the effect of applying the shift operators to the number 13:

```
unsigned short i, j;
```

```
i = 13;  
/* i is now 13 (binary 00000000000000001101) */
```

```
j = i << 2;  
/* j is now 52 (binary 000000000000110100) */
```

```
j = i >> 2;  
/* j is now 3 (binary 000000000000000011) */
```

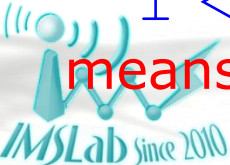
Bitwise Shift Operators (cont.)

- To modify a variable by shifting its bits, use the **compound assignment operators** `<<=` and `>>=`:

```
i = 13;  
/* i is now 13 (binary 000000000000001101) */  
  
i <<= 2;  
/* i is now 52 (binary 000000000000110100) */  
  
i >>= 2;  
/* i is now 13 (binary 000000000000001101) */
```

- The bitwise shift operators have lower precedence than the arithmetic operators, which can cause surprises:

`i << 2 + 1` not `(i << 2) + 1`
means `i << (2 + 1)`,

MSLab Since 2010

Bitwise Complement, *And*, Exclusive *Or*, and Inclusive *Or*

- There are four additional bitwise operators:

~ bitwise complement

& bitwise *and*

^ bitwise exclusive *or*

| bitwise inclusive *or*

| x | y | x y | x&y | x^y |
|---|---|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 |
| 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 1 | 1 | 0 |

- The ~ operator is unary; the other operators are binary.

Bitwise Complement, *And*, Exclusive *Or*, and Inclusive *Or* (cont.)

- Examples of the \sim , $\&$, \wedge , and \mid operators:

```
unsigned short i, j, k;
```

```
i = 21;
```

```
/* i is now      21 (binary 00000000000010101) */
```

```
j = 56;
```

```
/* j is now      56 (binary 00000000000111000) */
```

```
k = ~i;
```

```
/* k is now 65514 (binary 11111111111101010) */
```

```
k = i & j;
```

```
/* k is now      16 (binary 00000000000010000) */
```

```
k = i ^ j;
```

```
/* k is now      45 (binary 00000000000101101) */
```

```
k = i | j;
```

```
/* k is now      61 (binary 00000000000111101) */
```

Bitwise Complement, *And*, Exclusive *Or*, and Inclusive *Or* (cont.)

- The `~` operator can be used to help make low-level programs more portable.
- An integer whose bits are all 1: `~0`
- An integer whose bits are all 1 except for the last five: `~0x1f`

| | | | |
|----|-------------------|-------|-------------------|
| 0 | 00000000 00000000 | 0x1f | 00000000 00011111 |
| ~0 | 11111111 11111111 | ~0x1f | 11111111 11100000 |

Bitwise Complement, *And*, Exclusive *Or*, and Inclusive *Or* (cont.)

- Each of the \sim , $\&$, \wedge , and $|$ operators has a different precedence:

Highest: \sim

$\&$

\wedge

Lowest: $|$

- Examples:

$i \& \sim j | k$ means $(i \& (\sim j)) | k$

$i \wedge j \& \sim k$ means $i \wedge (j \& (\sim k))$

Using parentheses helps avoid confusion.


Bitwise Complement, *And*, Exclusive *Or*, and Inclusive *Or* (cont.)

- The **compound assignment** operators `&=`, `^=`, and `|=` correspond to the bitwise operators `&`, `^`, and `|`:

```
i = 21;  
/* i is now 21 (binary 00000000000010101) */  
  
j = 56;  
/* j is now 56 (binary 00000000000111000) */  
  
i &= j;  
/* i is now 16 (binary 00000000000010000) */  
  
i ^= j;  
/* i is now 40 (binary 00000000000101000) */  
  
i |= j;  
/* i is now 56 (binary 00000000000111000) */
```

Using the Bitwise Operators to Access Bits

- The bitwise operators can be used to **extract or modify** data stored in **a small number of bits**.
- Common single-bit operations:
 - Setting a bit
 - Clearing a bit
 - Testing a bit
- Assumptions:
 - `i` is a 16-bit unsigned short variable.
 - The leftmost—or ***most significant***—bit is numbered 15 and the least significant is numbered 0.



Using the Bitwise Operators to Access Bits (cont.)

- **Setting a bit.** The easiest way to set bit 4 of `i` is to *or* the value of `i` with the constant `0x0010`:

```
i = 0x0000;  
/* i is now 00000000000000000000 */
```

```
i |= 0x0010;  
/* i is now 000000000000000010000 */
```

- If the position of the bit is stored in the variable `j`, a *shift operator* can be used to create the mask:

```
i |= 1 << j;          /* sets bit j */
```

- Example: If `j` has the value 3, then `1 << j` is `0x0008`.



Using the Bitwise Operators to Access Bits (cont.)

- **Clearing a bit.** Clearing bit 4 of `i` requires a mask with a 0 bit in position 4 and 1 bits everywhere else:

```
i = 0x00ff;  
/* i is now 0000000011111111 */  
  
i &= ~0x0010;  
/* i is now 0000000011101111 */
```

- A statement that clears a bit whose position is stored in a variable:

```
i &= ~(1 << j); /* clears bit j */
```

`1 << j`

| | |
|----------|----------|
| 00000000 | 00001000 |
|----------|----------|

15 j 0

`~(1 << j)`

| | |
|----------|----------|
| 11111111 | 11110111 |
|----------|----------|

15 j 0


Using the Bitwise Operators to Access Bits (cont.)

- **Testing a bit.** An `if` statement that tests whether bit 4 of `i` is set:

```
if (i & 0x0010) ... /* tests bit 4 */
```

- A statement that tests whether bit `j` is set:

```
if (i & 1 << j) ... /* tests bit j */
```



Using the Bitwise Operators to Access Bits (cont.)

- Working with bits is easier if they are given names.
- Suppose that bits 0, 1, and 2 of a number correspond to the colors blue, green, and red, respectively.
- Names that represent the three bit positions:

```
#define BLUE 1
```

```
00000000 00000001
```

```
#define GREEN 2
```

```
00000000 00000010
```

```
#define RED 4
```

```
00000000 00000100
```

- Examples of setting, clearing, and testing the BLUE bit:

```
i |= BLUE;          /* sets BLUE bit */
```

```
i &= ~BLUE;         /* clears BLUE bit */
```

```
if (i & BLUE) ...   /* tests BLUE bit */
```

Using the Bitwise Operators to Access Bits (cont.)

- It's also easy to **set, clear, or test several bits at a time:**

```
i |= BLUE | GREEN;  
/* sets BLUE and GREEN bits */  
  
i &= ~(BLUE | GREEN);  
/* clears BLUE and GREEN bits */  
  
if (i & (BLUE | GREEN)) ...  
/* tests BLUE and GREEN bits */
```

BLUE | GREEN
00000000 00000011

- The `if` statement tests **whether either the BLUE bit or the GREEN bit is set.**

Using the Bitwise Operators to Access Bit-Fields

- Dealing with a group of several consecutive bits (a **bit-field**) is slightly more complicated than working with single bits.
- Common bit-field operations:
 - Modifying a bit-field
 - Retrieving a bit-field

Using the Bitwise Operators to Access Bit-Fields (cont.)

- **Modifying a bit-field.** Modifying a bit-field requires two operations:
 - A bitwise *and* (to clear the bit-field)
 - A bitwise *or* (to store new bits in the bit-field)
- Example:

```
i = i & ~0x0070 | 0x0050;
/* stores 101 in bits 4-6 */
```

$\sim 0x0070$

11111111 10001111

$0x0050$

00000000 01010000

 - The `&` operator clears bits 4–6 of `i`; the `|` operator then sets bits 6 and 4.

Using the Bitwise Operators to Access Bit-Fields (cont.)

- To generalize the example, assume that *j* contains the value to be stored in bits 4–6 of *i*.
- *j* will need to be shifted into position before the bitwise *or* is performed:

```
i = (i & ~0x0070) | (j << 4);
```

```
/* stores j in bits 4-6 */ j << 4
```

j

00000000 00000**101**

j << 4

00000000 0**101**0000

Using the Bitwise Operators to Access Bit-Fields (cont.)

- **Retrieving a bit-field.** Fetching a bit-field at the right end of a number (in the least significant bits) is easy:

```
j = i & 0x0007;  
/* retrieves bits 0-2 */
```

0x0007 00000000 00000**111**

- If the bit-field **isn't at the right end of i**, we can **first shift the bit-field to the end** before extracting the field using the & operator:

```
j = (i >> 4) & 0x0007;  
/* retrieves bits 4-6 */
```

i **???????? ?xxx????**

i >> 4 **???????? ?????xxx**

Program: XOR Encryption

- One of the simplest ways to encrypt data is to exclusive-or (XOR) each character with a secret key.
- Suppose that the key is the & character.
- XORing this key with the character z yields the \ character:

| | | |
|-----|-----------------|--------------------|
| | 00100110 | (ASCII code for &) |
| XOR | <u>01111010</u> | (ASCII code for z) |
| | 01011100 | (ASCII code for \) |

Program: XOR Encryption (cont.)

- Decrypting a message is done by applying the same algorithm:

| | | |
|-----|-----------------|--------------------|
| | 00100110 | (ASCII code for &) |
| XOR | 01011100 | (ASCII code for \) |
| | <u>01111010</u> | (ASCII code for z) |

Program: XOR Encryption (cont.)

- The `xor.c` program encrypts a message by XORing each character with the `&` character.
- The original message can be entered by the user or read from a file using input redirection.
- The encrypted message can be viewed on the screen or saved in a file using output redirection.
- A sample file named `msg`:

Trust not him with your secrets, who, when left alone in your room, turns over your papers.

--Johann Kaspar Lavater (1741-1801)

Program: XOR Encryption (cont.)

- A command that **encrypts** `msg`, **saving** the encrypted message **in** `newmsg`:

```
xor <msg >newmsg
```

- Contents of `newmsg`:

```
rTSUR HIR NOK QORN _IST UCETCRU, QNI, QNCH JC@R  
GJIHC OH _IST TIIK, RSTHU IPCT _IST VGVCTU.  
--LINGHH mGUVGT jGPGRCT (1741-1801)
```

- A command that **recovers** the original message and displays it on the screen:

```
xor <newmsg
```

Program: XOR Encryption (cont.)

- The `xor.c` program **won't change some characters**, including digits.
- **XORing these characters with & would produce invisible control characters**, which could cause problems with some operating systems.
- The program **checks whether** both the original character and the new (encrypted) character are **printing characters**.
- **If not**, the program will **write the original character** instead of the new character.

Program: XOR Encryption (cont.)

xor.c

```
#include <ctype.h>
#include <stdio.h>

#define KEY '&'

int main(void)
{
    int orig_char, new_char;

    while ((orig_char = getchar()) != EOF) {
        new_char = orig_char ^ KEY;
        if (isprint(orig_char) && isprint(new_char))
            putchar(new_char);
        else
            putchar(orig_char);
    }

    return 0;
}
```